

```
import time
from collections import deque

import numpy as np
import tensorflow as tf
import tensorflow.keras as keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import InputLayer, Dense

class DQN:
    '''
        Baseline Deep Q-Network with experience replay
    '''

    def __init__(self, state_size, action_size, policy, learning_delay, loss_fn, epsilon,
gamma,
        learning_rate, n_units, buffer_size, l2=0, learning_freq=1, verbose=False, **kwar
gs):
        '''
            Initialize necessary fields
        '''
        self.type = "DQN"

        self.state_size = state_size
        self.action_size = action_size
        self.policy = policy
        self.learning_delay = learning_delay
        self.learning_freq = learning_freq
        self.loss_fn = loss_fn
        self.epsilon = epsilon
        self.gamma = gamma
        self.optimizer = keras.optimizers.Adam(lr=learning_rate)
        self.setup_model(n_units, l2=l2)

        self.epsilon_log = []
        self.reward_log = []
        self.loss_log = []
        self.deque_log = []
        self.verbose = verbose
        self.replay_buffer = {
            "states": deque(maxlen=buffer_size),
            "actions": deque(maxlen=buffer_size),
            "rewards": deque(maxlen=buffer_size),
            "next_states": deque(maxlen=buffer_size),
            "done": deque(maxlen=buffer_size)
        }
        self.episode = 0

    def build_model(self, n_units, activation="elu", l2=0):
        '''
            Build a simple sequential model.
        '''
        print("L2: {}".format(l2))

        model = Sequential()
        i = 0

        # Input layer
        model.add(InputLayer(input_shape=(self.state_size,)))

        # Loop over hidden layers
        for n in n_units:
            model.add(Dense(n,
                activation=activation,
                kernel_regularizer=keras.regularizers.l2(l2),
                name = "D"+str(i)))
            i=i+1

        # model.add(BatchNormalization())
```

```

    # Output layer
    model.add(Dense(self.action_size,
                    activation=None,
                    name = "D"+str(i)))

    return model

def setup_model(self, n_units, l2=0):
    """
    Compile a simple sequential model
    """

    model = self.build_model(n_units=n_units, l2=l2)
    self.model = model
    model.summary()

def get_epsilon(self):
    try:
        return self.epsilon(self.episode)
    except TypeError as e:
        return self.epsilon

def play_one_step(self, env, state):
    """
    Take one step in the environment based on the agent parameters
    """

    action = self.policy(state, self.model, self.get_epsilon()) # Query policy

    next_state, reward, done, info = env.step(action) # Query environment
    self.memorize(state, action, reward, next_state, done) # Log

    return next_state, reward, done, info

def memorize(self, state, action, reward, next_state, done):
    """
    Log the experience from one step into the replay buffer as a dictionary
    """

    state = np.array(state, ndmin=2)
    next_state = np.array(next_state, ndmin=2)

    self.replay_buffer["states"].append(state)
    self.replay_buffer["actions"].append(action)
    self.replay_buffer["rewards"].append(reward)
    self.replay_buffer["next_states"].append(next_state)
    self.replay_buffer["dones"].append(done)

def sample_experience_inds(self, batch_size):
    """
    Sample batch_size number of experience indices from the replay buffer
    """

    # If batch size greater than current length of buffer, give all indices for buffer.
    # Otherwise, get random sampling of batch_size indices.
    choice_range = len(self.replay_buffer["states"])
    if batch_size is None or batch_size > choice_range:
        indices = np.random.choice(choice_range, size=choice_range, replace=False)
    else:
        indices = np.random.choice(choice_range, size=batch_size, replace=False)

    return indices

def sample_experience_inds_old(self, batch_size):
    """
    Sample batch_size number of experience indices from the replay buffer
    """

    # If batch size greater than current length of buffer, give all indices for buffer

```

```

r.
    # Otherwise, get random sampling of batch_size indices.
    if batch_size > len(self.replay_buffer["states"]):
        indices = list(range(len(self.replay_buffer["states"])))
    else:
        indices = np.random.randint(len(self.replay_buffer["states"]), size=batch_size)

e)

    return indices

def sample_experience(self, inds):
    '''
    Sample experiences with indices from replay buffer
    '''

    batch = {}
    for key in self.replay_buffer.keys():
        batch[key] = [self.replay_buffer[key][index] for index in inds]

    batch["states"] = np.concatenate(batch["states"], axis=0)
    batch["next_states"] = np.concatenate(batch["next_states"], axis=0)

    return batch

def get_current_Q_values(self, states):
    return self.model(states)

def get_next_Q_values(self, next_states):
    return self.model.predict(next_states)

def learning_step(self, batch_size=100):
    '''
    Train the model with one batch by sampling from replay buffer
    Use the gradient tape method
    '''

    batch_time_start = time.time()
    # Fetch batch
    batch_inds = self.sample_experience_inds(batch_size)
    batch = self.sample_experience(batch_inds)

    # Create target q values, with mask to disregard irrelevant actions
    next_Q_values = self.get_next_Q_values(batch["next_states"]) # Get subsequent Q values
    max_next_Q_values = np.max(next_Q_values, axis=1) # Get max of subsequent Q values
    target_Q_values = (batch["rewards"] + (1 - np.asarray(batch["dones"]))) * self.gamma * max_next_Q_values # Define Q targets
    mask = tf.one_hot(batch["actions"], self.action_size) # Create mask to mask actions not taken

    # Use optimizer to apply gradient to model
    with tf.GradientTape() as tape:
        all_Q_values = self.get_current_Q_values(batch["states"]) # Get all possible q values from the states
        masked_Q_values = all_Q_values * mask # Mask the actions which were not taken
        Q_values = tf.reduce_sum(masked_Q_values, axis=1) # Get the sum to reduce to action taken
        loss = tf.reduce_mean(self.loss_fn(target_Q_values, Q_values)) # Compute the losses
        self.loss_log.append(loss) # Append to log
        grads = tape.gradient(loss, self.model.trainable_variables) # Compute the gradients
        self.optimizer.apply_gradients(zip(grads, self.model.trainable_variables)) # Apply the gradients to the model

    def execute_episode(self, env, n_steps=None, render_flag=False, batch_size=100, verbose=False, train=True):

```

```

'''
Execute one episode, which terminates when done if flagged or step limit is reach
ed
'''

# Initialize vars
reward_total = 0
step = 0
done = False
state = env.reset()
while (n_steps is None or step < n_steps) and not done: # Continue till step coun
t, or until done

    if render_flag: # Create visualization for environment
        env.render()

    state, reward, done, info = self.play_one_step(env, state) # Custom step func
tion

    reward_total += reward
    step += 1
    if done:
        break

# If train flag and episode above some threshold (to fill buffer), train
if train and self.episode >= self.learning_delay and self.episode % self.learning
_freq == 0:
    print("\tLearning")
    self.learning_step(batch_size=batch_size)
else:
    print("\tCollecting")

self.reward_log.append(reward_total)
self.epsilon_log.append(self.get_epsilon())
self.deque_log.append(len(self.replay_buffer["states"]))
self.episode += 1

if verbose:
    print("\tReward: {}".format(reward_total))

def execute_episodes(self, env, n_episodes, n_steps, render_flag=False, batch_size=10
0, verbose=False,
train=True):
'''
Execute multiple episodes
'''

for episode in range(n_episodes):
    if verbose:
        print("Episode: {}".format(self.episode))

    self.execute_episode(
        env=env,
        n_steps=n_steps,
        render_flag=render_flag,
        batch_size=batch_size,
        verbose=verbose,
        train=train)

    if render_flag:
        env.close()

```